

CSDS 440: Machine Learning

Soumya Ray (he/him, sray@case.edu)

Olin 516

Office hours T, Th 11:15-11:45 or by appointment

[Zoom Link](#)

Announcements

- Quiz 2 Thursday
 - Same format as Q1
 - Topics: Everything up to and including Lecture 8 (Calculus)
 - Bring a scientific calculator (or phone app)

Recap

- Having very few layers is problematic because they will need many n_____ and not have s_____.
- A further motivation is from C_____ c_____. This creates a deep architecture by successively learning p_____, each modeling the r_____ of the previous.
- One way to interpret hidden units is as f_____ c_____ for an h____ d_____ space where classification can be done with a p_____.
- To generate good features, we should allow l_____ c_____ p_____.
- We should also a_____ i_____ across different parts of the input.
- Neural Networks can be viewed as c_____ g_____. Each layer of a network performs a m_____ operation on the input v_____.
- Fully connected layers can be problematic because the number of weights g_____ q_____.
- We should let each hidden unit only look at a l_____ part of the input.
- If the weights connecting a region to hidden units are the same across hidden units, this is called i_____.
- In a c_____ neural network, we introduce a k_____. This is a set of weights that are r_____ across multiple l_____ regions.

Today

- Artificial Neural Networks (Ch 4, Mitchell)
- Probabilistic Machine Learning

How to scale an ANN?

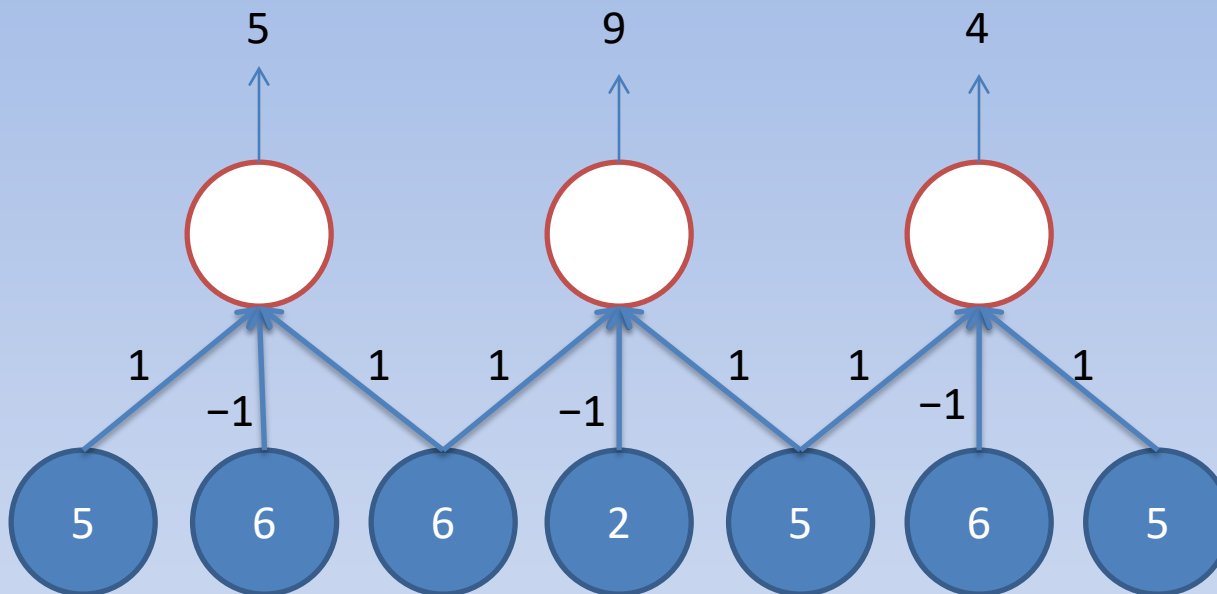
Suppose we create an ANN with LOTS of layers.

- ~~1. Why might we want to do that?~~
2. What will these layers *do*?
3. How can learning scale?
4. How to deal with vanishing gradients?
5. How to deal with overfitting?

Convolutional Neural Networks

- Introduce a **kernel** k : a set of weights *replicated* across multiple local regions
 - Generally multiple such kernels will be used
 - Each kernel computes one local feature
- The operation of applying the kernel to the input is called **convolution**

Convolution



$k=[1,-1,1]$, size $l=3$, stride $s=2$

Convolution

$$\mathbf{z} = \mathbf{x} * \mathbf{k}$$

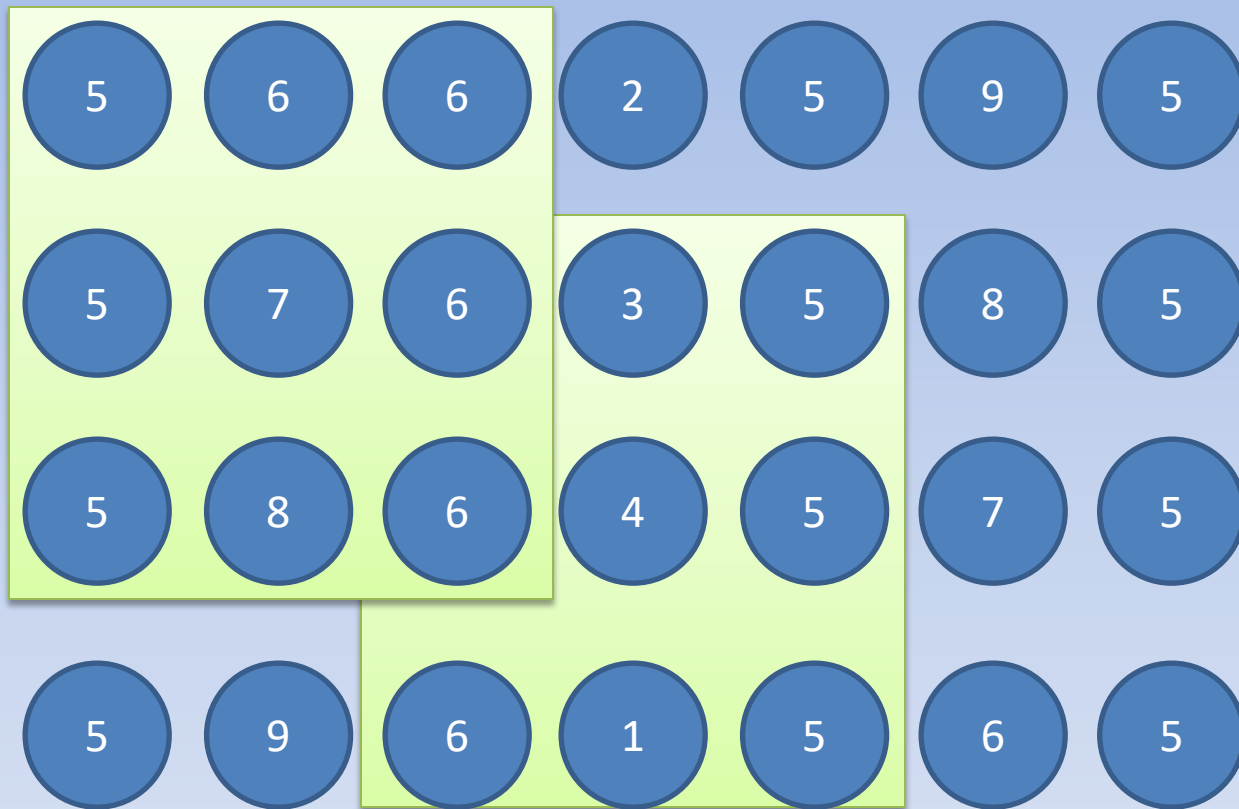
$$z_i = \sum_{j=1}^l k_j x_{i+j-(l+1)/2}$$

- Convolution is a linear operation
 - Can also be represented as a matrix operation
- The output \mathbf{z} will have roughly n/s entries

Convolutional NNs

- A kernel detects a specific feature, but what kernels to use?
- In a CNN, the kernels (detectors/filters) *themselves can be learned*
 - Parameterize as a set of weights, and learn via backpropagation

2D Convolution

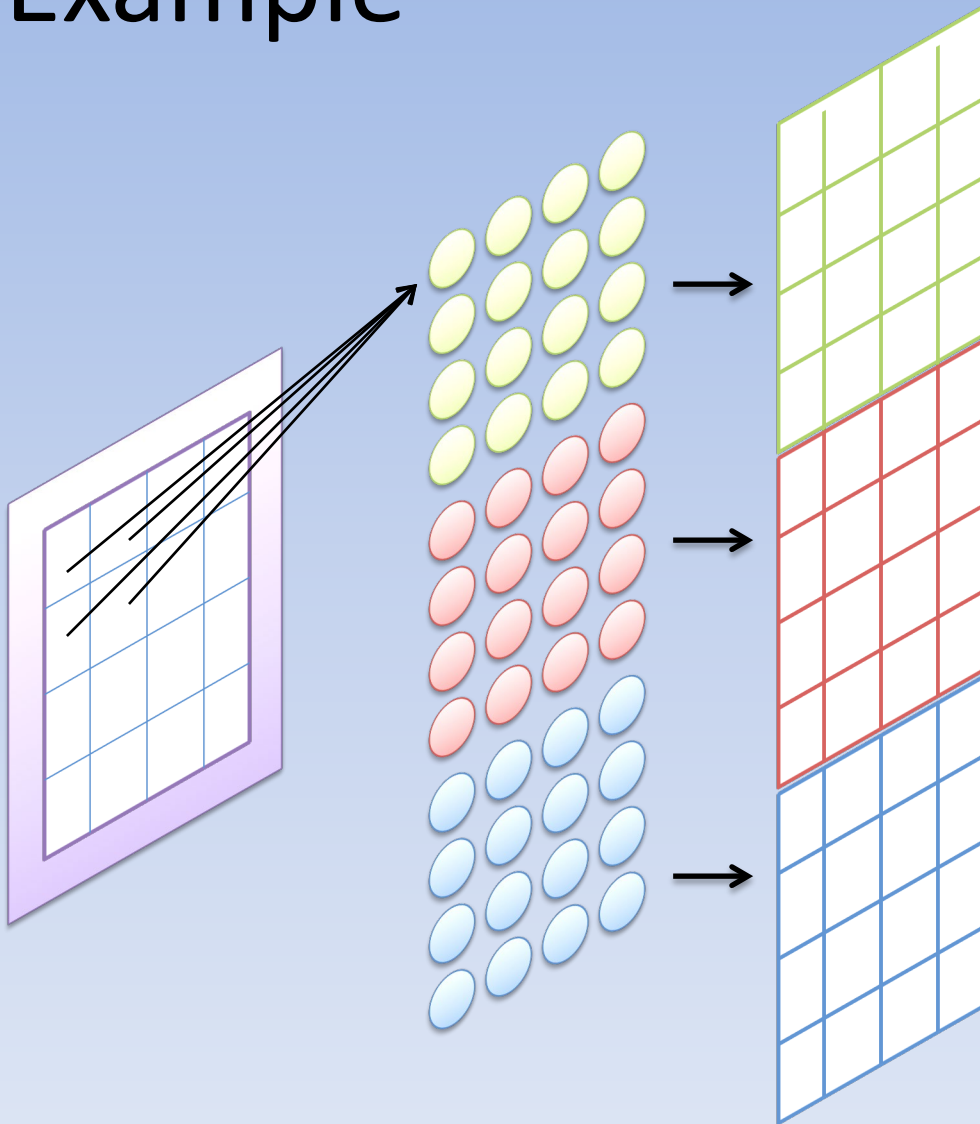


$$k=[w_1 \ w_2 \ w_3; w_4 \ w_5 \ w_6; w_7 \ w_8 \ w_9]; \text{ stride}=(2,1)$$

Tensors

- Each convolution kernel creates one “feature detector” that is looking for a specific property in a local patch
- Typically, will have many such kernels, each looking for a different feature

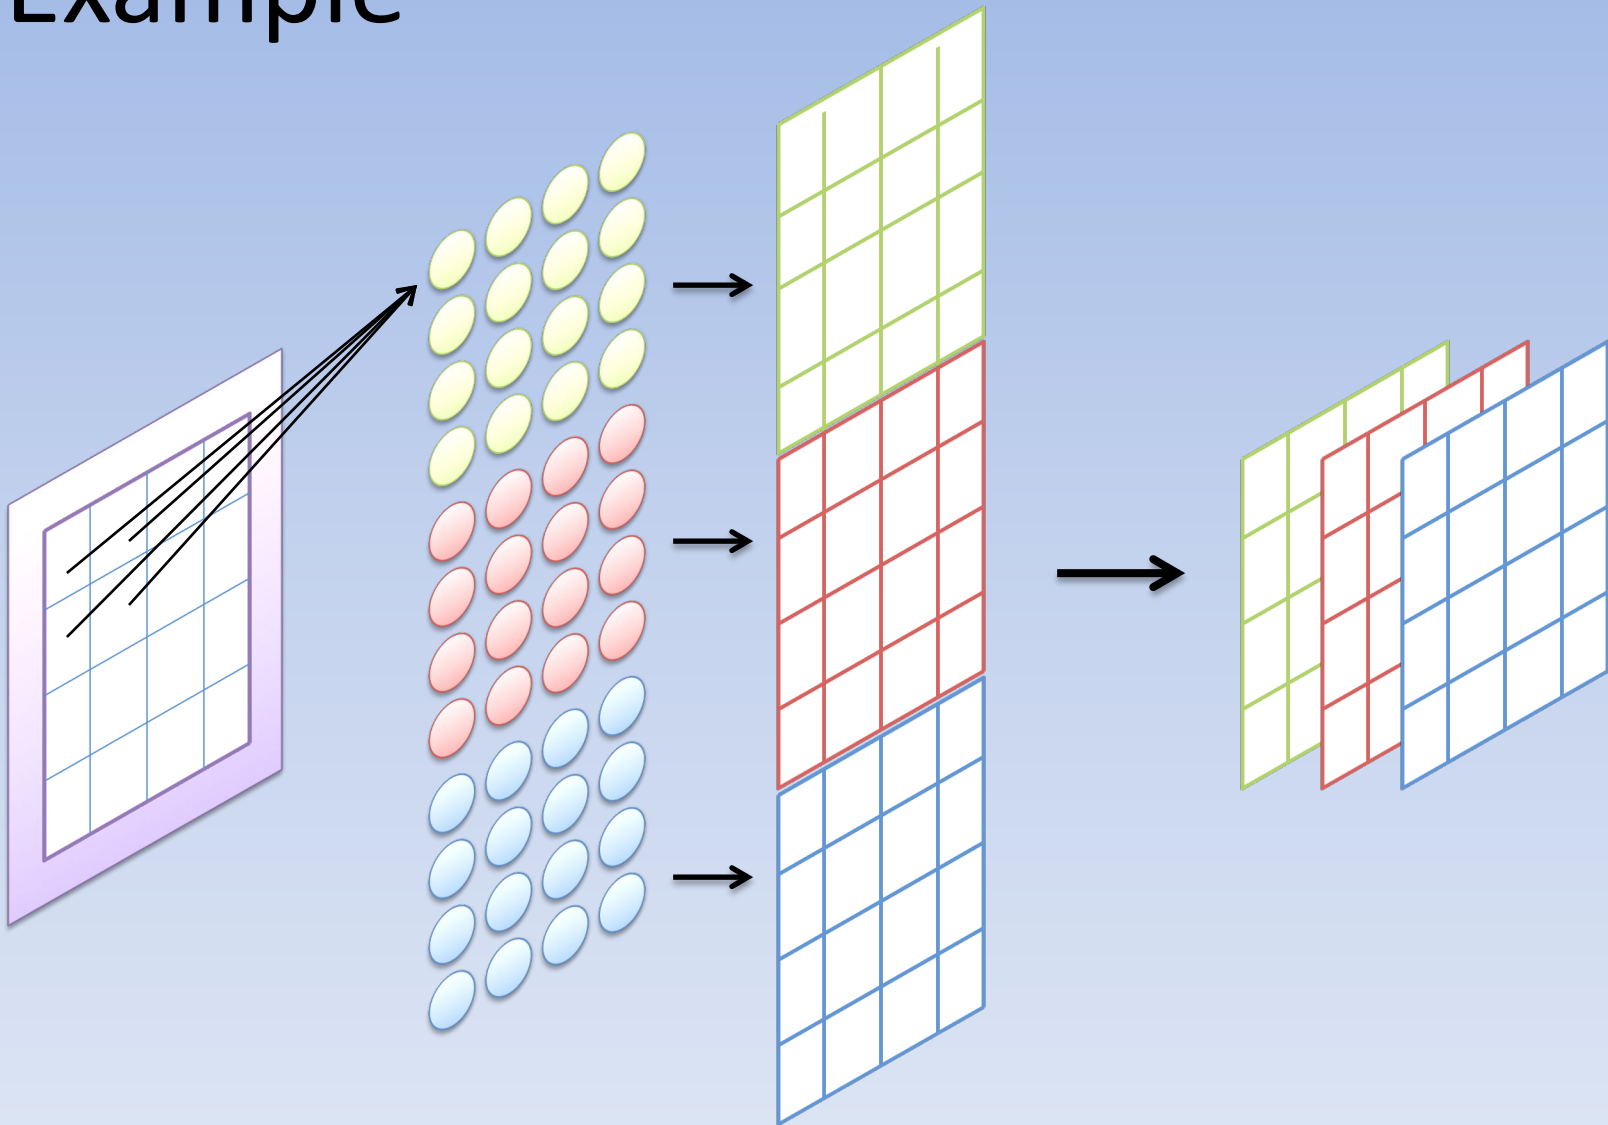
Example



Tensors

- In order to maintain locality and invariance, instead of concatenating these kernels, we stack them along a new dimension
- If the input has two dimensions or more (e.g. images, video) then this results in a multidimensional matrix at each layer
 - These create “**tensors**”

Example



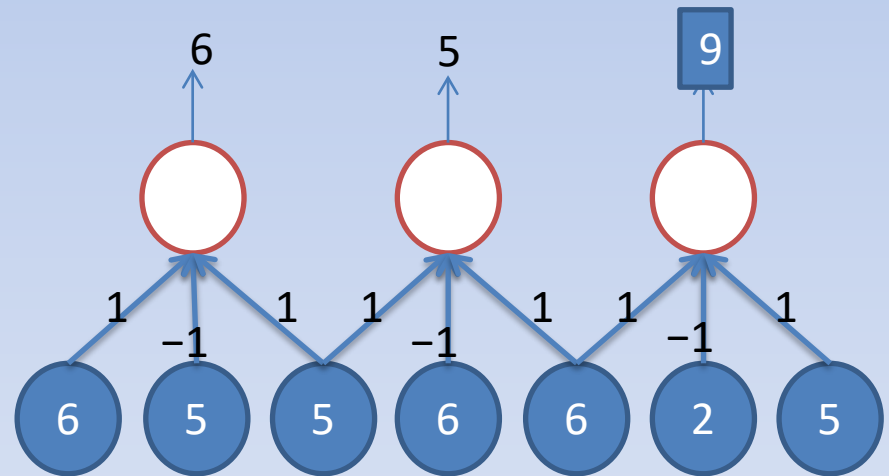
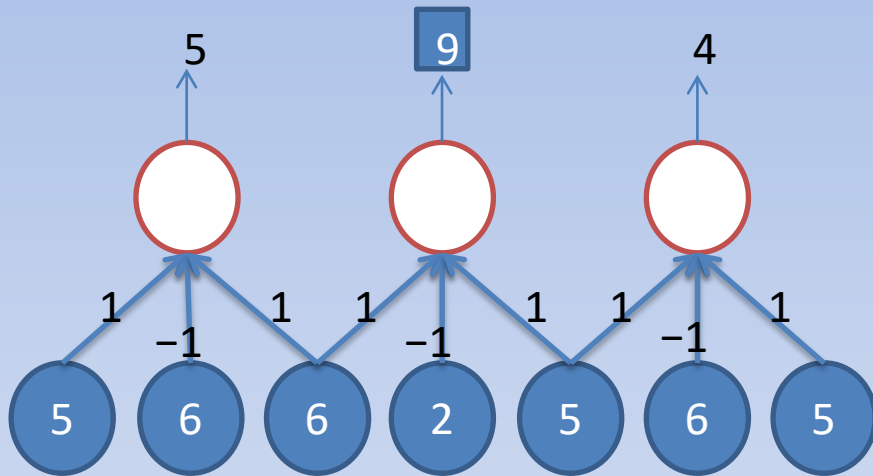
Example

- Suppose we have a batch of 32x32 images
 - Each input has dimensions (32, 32)
- Suppose we apply 10 4x4 kernels to each image with stride 1x1
- The output will be a tensor with dimensions (10, 32, 32)

Pooling Layers

- A *pooling* layer aggregates information from an adjacent layer
- Average pooling: $k=(1/l, 1/l, \dots 1/l)$
- Max pooling: computes the maximum value of l inputs
 - For each feature detector, identifies whether that feature was found somewhere in the previous layer
- Downsamples input by factor of l

Max Pooling



Vanishing Gradients 1

- A key problem in ANNs is *vanishing gradients*
- To prevent vanishing gradients, we can use the “Rectified Linear Unit” (ReLU) activation function:

$$h(x) = \max(0, x)$$

Vanishing Gradients 2

- Each layer in an ANN *learns a completely new representation* from the previous layer
 - Can cause catastrophic failure due to one “bad” layer
- Instead, each layer can *add on to* the learned representation of the previous layer
 - Allows building much deeper structures robustly

Residual Networks

- Perturbing the representation is done through adding a “residual” function to each layer
- Replace

$$\mathbf{z}^{l+1} = h(\mathbf{W}^l \mathbf{z}^l)$$

- With

$$\mathbf{z}^{l+1} = h(\mathbf{z}^l + f(\mathbf{z}^l))$$

Residual function
(learned from
data; could be
identity)

How do we prevent overfitting?

- ANNs are very prone to overfitting
 - Structure can be very complex, lots of parameters
 - Decision surface can be very nonlinear

Controlling Overfitting

- One strategy: add a “weight decay” term

$$L_{OC}(\mathbf{w}) = L(\mathbf{w}) + \gamma \sum_i \sum_j w_{ji}^2$$

Weight Decay
Term
“Complexity
Penalty”

Tradeoff Hyperparameter

- This will prevent weights from growing too large